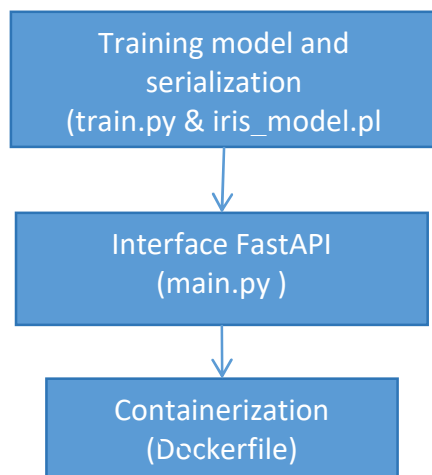
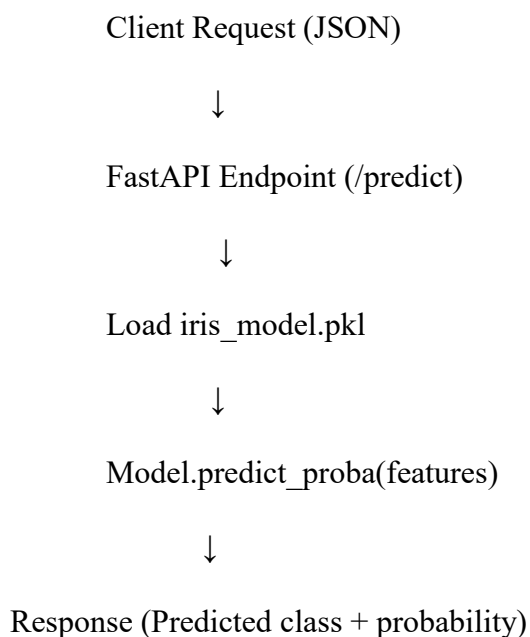


Project Title: Database & Python ETL with Reproducible Infrastructure

1. **Objective :** To deploy a Machine Learning model using FastAPI and Docker, providing a lightweight and reproducible setup for serving ML predictions through REST API endpoints. The project demonstrates how to load a trained model and deploy as an API endpoint fully reproducible via Docker.
2. **Architecture Overview :**



Data Flow :



3. Technology Stack :

Component	Technology	Purpose
Programming Language	Python 3.10	ML logic & API framework
Framework	FastAPI	RestfulAPI for prediction
Libraries	Scikit-learn,pydantic,numpy	Model training & inference
Serialization	Pickle	ML saved model
Web server	Uvicorn	Server for fastAPI
Containerization	Docker	Reproducible Deployment

4. Data Flow

Model Training (train.py)

- A set of numeric features (16 parameters) representing patient attributes relevant to thyroid detection.
- Loads the Iris dataset from sklearn.datasets.
- Constructs a Pipeline:
 - ♦ StandardScaler for normalization
 - ♦ LogisticRegression for classification
- Splits data into train/test sets (80/20).
- Trains the pipeline and evaluates accuracy (~95–98%).
- Serializes the trained model into models/iris_model.pkl using pickle

Interface(main.py)

- Loads the model (iris_model.pkl) is loaded via the pickle module.

```
22 def load_model():
23     if not MODEL_PATH.exists():
24         raise FileNotFoundError(f"Model file not found at {MODEL_PATH.resolve()}. "
25                                 f"Run `python train.py` first.")
26     with open(MODEL_PATH, "rb") as f:
27         model = pickle.load(f)
28     return model
```

- Validates incoming JSON requests with Pydantic.

```
from pydantic import BaseModel, Field, conlist

IrisFeatureVector = conlist(float, min_length=4, max_length=4)
```

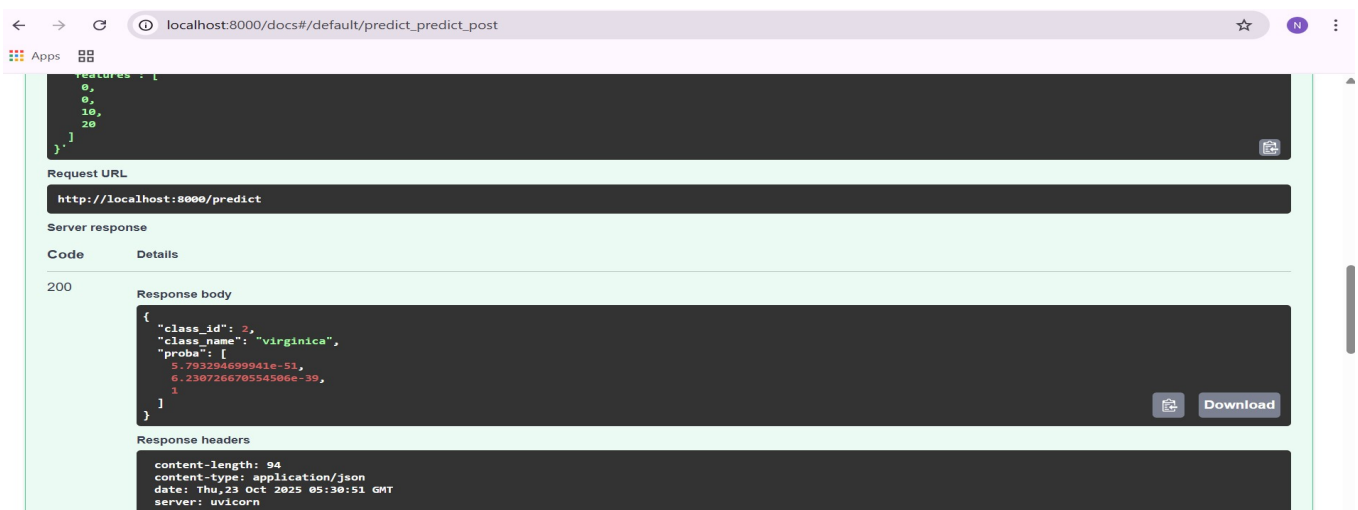
- Passes the feature vector to model.predict_proba().

```
@app.post("/predict", response_model=PredictResponse)
def predict(req: PredictRequest):
    try:
        X = [req.features]
        proba = model.predict_proba(X)[0].tolist()
        cls = int(max(range(len(proba)), key=lambda i: proba[i]))
        return PredictResponse(class_id=cls, class_name=CLASS_MAP[cls], proba=proba)
    except Exception as e:
        raise HTTPException(status_code=400, detail=str(e))
```

- Returns the predicted class name (setosa, versicolor, or virginica) with probabilities.

Output

- The ML model is deployed on an API and runs successfully with prediction



localhost:8000/docs#/default/predict_predict_post

Request URL

http://localhost:8000/predict

Server response

Code Details

200

Response body

```
{
  "class_id": 2,
  "class_name": "virginica",
  "proba": [
    5.793294699941e-51,
    6.230726670554506e-39,
    1
  ]
}
```

Response headers

```
content-length: 94
content-type: application/json
date: Thu, 23 Oct 2025 05:30:51 GMT
server: uvicorn
```

5. API Specification

/health : get() - Health check and returns the status

/predict : post() - Takes 4 features as input and returns prediction with probability

6. Environmental Variables

MODEL_PATH : Path to serialized model - iris_model.pkl

PORT : FastAPI Port : 8000

7. Dockerisation

Container - iris_api

```
Dockerfile > ...
1 FROM python:3.11-slim
2
3 ENV PYTHONUNBUFFERED=1 \
4     PYTHONDONTWRITEBYTECODE=1
5
6 WORKDIR /app
7
8 RUN apt-get update && apt-get install -y --no-install-recommends \
9     build-essential \
10    && rm -rf /var/lib/apt/lists/*
11
12 COPY requirements.txt .
13 RUN pip install --no-cache-dir -r requirements.txt
14
15 COPY app/ app/
16
17 COPY iris_model.pkl .
18
19 EXPOSE 8000
20 CMD ["uvicorn", "app.main:app", "--host", "0.0.0.0", "--port", "8000"]
21
```

8. Testing & Validation

Launch - Check Health : Returns Status: ok

The screenshot shows a web browser window with the address bar displaying `localhost:8000/docs#/default/health_health_get`. Below the browser window is a REST client interface. The 'Curl' section shows the command: `curl -X 'GET' \ 'http://localhost:8000/health' \ -H 'accept: application/json'`. The 'Request URL' section shows `http://localhost:8000/health`. The 'Server response' section shows a 200 status code. The 'Response body' section shows a JSON object: `{ "status": "ok" }`. The 'Response headers' section shows: `content-length: 15, content-type: application/json, date: Thu, 23 Oct 2025 05:56:49 GMT, server: uvicorn`.

Prediction : localhost:8000/predict : Predict which class it belongs.

POST /predict Predict

Parameters

No parameters

Request body required

application/json

Edit Value | Schema

```
{  "features": [    0,    0,   10,   20  ]}
```

200

Response body

```
{  "class_id": 2,  "class_name": "virginica",  "proba": [    5.793294699941e-51,    6.230726670554506e-39,    1  ]}
```

Response headers

```
content-length: 94
content-type: application/json
date: Thu, 23 Oct 2025 05:30:51 GMT
server: uvicorn
```

Validate Swagger UI at → <http://localhost:8000/docs>

9. Reproducibility

- All components (model, code, dependencies) are defined under Docker.
- requirements.txt fully pins package versions
- The solution can be build using the commands
 - Docker build --no cache -t iris-api
 - Git : [GitHub](#)